

## Import Libraries

In [1]: `pip install yellowbrick`

```
Requirement already satisfied: yellowbrick in c:\users\lenovo\anaconda3\lib\site-packages (1.5)
Requirement already satisfied: cyclopy>=0.10.0 in c:\users\lenovo\anaconda3\lib\site-packages (from
yellowbrick) (0.11.0)
Requirement already satisfied: scikit-learn>=1.0.0 in c:\users\lenovo\anaconda3\lib\site-packages
(from yellowbrick) (1.0.2)
Requirement already satisfied: numpy>=1.16.0 in c:\users\lenovo\anaconda3\lib\site-packages (from
yellowbrick) (1.21.5)
Requirement already satisfied: matplotlib!=3.0.0,>=2.0.2 in c:\users\lenovo\anaconda3\lib\site-pac
kages (from yellowbrick) (3.5.2)
Requirement already satisfied: scipy>=1.0.0 in c:\users\lenovo\anaconda3\lib\site-packages (from y
ellowbrick) (1.9.1)
Requirement already satisfied: pyparsing>=2.2.1 in c:\users\lenovo\anaconda3\lib\site-packages (fr
om matplotlib!=3.0.0,>=2.0.2->yellowbrick) (3.0.9)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\lenovo\anaconda3\lib\site-packages (f
rom matplotlib!=3.0.0,>=2.0.2->yellowbrick) (4.25.0)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\lenovo\anaconda3\lib\site-packages
(from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (2.8.2)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\lenovo\anaconda3\lib\site-packages (f
rom matplotlib!=3.0.0,>=2.0.2->yellowbrick) (1.4.2)
Requirement already satisfied: packaging>=20.0 in c:\users\lenovo\anaconda3\lib\site-packages (fro
m matplotlib!=3.0.0,>=2.0.2->yellowbrick) (21.3)
Requirement already satisfied: pillow>=6.2.0 in c:\users\lenovo\anaconda3\lib\site-packages (from
matplotlib!=3.0.0,>=2.0.2->yellowbrick) (9.2.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\lenovo\anaconda3\lib\site-packages
(from scikit-learn>=1.0.0->yellowbrick) (2.2.0)
Requirement already satisfied: joblib>=0.11 in c:\users\lenovo\anaconda3\lib\site-packages (from s
cikit-learn>=1.0.0->yellowbrick) (1.1.0)
Requirement already satisfied: six>=1.5 in c:\users\lenovo\anaconda3\lib\site-packages (from pytho
n-dateutil>=2.7->matplotlib!=3.0.0,>=2.0.2->yellowbrick) (1.16.0)
Note: you may need to restart the kernel to use updated packages.
```

```
WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)
```

In [2]:

```
import numpy as np
import pandas as pd
from pandas.plotting import parallel_coordinates

import os
import sqlite3
import math
from collections import Counter
from pathlib import Path
from tqdm import tqdm

import seaborn as sns
import matplotlib as mpl
import matplotlib.pyplot as plt
import plotly
import plotly.graph_objects as go
import plotly.express as px
from plotly.subplots import make_subplots
import plotly.io as pio

from scipy.stats import skew
import yellowbrick
import sklearn
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.manifold import TSNE
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.metrics import roc_auc_score
```

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV

mpl.rcParams['font.family'] = 'monospace'
sns.set_theme(style="white", palette=None)
plotly.offline.init_notebook_mode()
plt.rcParams['figure.dpi'] = 300
plt.rcParams['savefig.dpi'] = 300
```


In [3]: `%matplotlib inline`

In [4]: `# Reading csv files and drop the first column`  
`df_train = pd.read_csv("fraudTrain.csv")`  
`df_train.drop(df_train.columns[0], axis=1, inplace=True)`  
  
`df_test = pd.read_csv("fraudTest.csv")`  
`df_test.drop(df_test.columns[0], axis=1, inplace=True)`  
  
`# First view 10 rows`  
`df_train.head(10)`

Out[4]:

	trans_date_trans_time	cc_num	merchant	category	amt	first	last	gender	street
0	2019-01-01 00:00:18	2703186189652095	fraud_Rippin, Kub and Mann	misc_net	4.97	Jennifer	Banks	F	56' Perry Cove
1	2019-01-01 00:00:44	630423337322	fraud_Heller, Gutmann and Zieme	grocery_pos	107.23	Stephanie	Gill	F	4303' Riley Green Suite 39'
2	2019-01-01 00:00:51	38859492057661	fraud_Lind-Buckridge	entertainment	220.11	Edward	Sanchez	M	59' White Dale Suite 53'
3	2019-01-01 00:01:16	3534093764340240	fraud_Kutch, Hermiston and Farrell	gas_transport	45.00	Jeremy	White	M	944' Cynthia Cour Apt 03'
4	2019-01-01 00:03:06	375534208663984	fraud_Keeling-Crist	misc_pos	41.96	Tyler	Garcia	M	40' Bradley Res
5	2019-01-01 00:04:08	4767265376804500	fraud_Stroman, Hudson and Erdman	gas_transport	94.63	Jennifer	Conner	F	465' David Island
6	2019-01-01 00:04:42	30074693890476	fraud_Rowe-Vandervort	grocery_net	44.54	Kelsey	Richards	F	88' Sarah Station Suite 62'
7	2019-01-01 00:05:08	6011360759745864	fraud_Corwin-Collins	gas_transport	71.65	Steven	Williams	M	23' Flore Pas Suite 72'
8	2019-01-01 00:05:18	4922710831011201	fraud_Herzog Ltd	misc_pos	4.27	Heather	Chase	F	688' Hick Streat Suite 95'
9	2019-01-01 00:06:01	2720830304681674	fraud_Schoen, Kuphal and Nitzsche	grocery_pos	198.39	Melissa	Aguilar	F	2132' Taylo Square Suite 70'

10 rows × 22 columns



```
In [5]: df_train.columns
```

```
Out[5]: Index(['trans_date_trans_time', 'cc_num', 'merchant', 'category', 'amt', 'first', 'last', 'gender', 'street', 'city', 'state', 'zip', 'lat', 'long', 'city_pop', 'job', 'dob', 'trans_num', 'unix_time', 'merch_lat', 'merch_long', 'is_fraud'], dtype='object')
```

```
In [6]: #Rename Columns
df_train.rename(columns={"trans_date_trans_time": "transaction_time", "cc_num": "credit_card_number",
```

```
In [7]: #convert date time column
df_train["transaction_time"] = pd.to_datetime(df_train["transaction_time"], infer_datetime_format=True)
df_train["dob"] = pd.to_datetime(df_train["dob"], infer_datetime_format=True)
```

```
In [8]: from datetime import datetime

# Apply function utcfromtimestamp and drop column unix_time
df_train['time'] = df_train['unix_time'].apply(datetime.utcfromtimestamp)
df_train.drop('unix_time', axis=1)

# Add cloumn hour of day
df_train['hour_of_day'] = df_train.time.dt.hour
```

```
In [9]: df_train[['time', 'hour_of_day']]
```

```
Out[9]:
```

	time	hour_of_day
0	2012-01-01 00:00:18	0
1	2012-01-01 00:00:44	0
2	2012-01-01 00:00:51	0
3	2012-01-01 00:01:16	0
4	2012-01-01 00:03:06	0
...	...	...
1296670	2013-06-21 12:12:08	12
1296671	2013-06-21 12:12:19	12
1296672	2013-06-21 12:12:32	12
1296673	2013-06-21 12:13:36	12
1296674	2013-06-21 12:13:37	12

1296675 rows × 2 columns

```
In [10]: df_train.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1296675 entries, 0 to 1296674
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   transaction_time      1296675 non-null  datetime64[ns]
1   credit_card_number    1296675 non-null  int64
2   merchant              1296675 non-null  object
3   category              1296675 non-null  object
4   amount(usd)          1296675 non-null  float64
5   first                 1296675 non-null  object
6   last                  1296675 non-null  object
7   gender                1296675 non-null  object
8   street                1296675 non-null  object
9   city                  1296675 non-null  object
10  state                 1296675 non-null  object
11  zip                   1296675 non-null  int64
12  lat                   1296675 non-null  float64
13  long                  1296675 non-null  float64
14  city_pop              1296675 non-null  int64
15  job                   1296675 non-null  object
16  dob                   1296675 non-null  datetime64[ns]
17  transaction_id        1296675 non-null  object
18  unix_time             1296675 non-null  int64
19  merch_lat             1296675 non-null  float64
20  merch_long            1296675 non-null  float64
21  is_fraud              1296675 non-null  int64
22  time                  1296675 non-null  datetime64[ns]
23  hour_of_day           1296675 non-null  int64
dtypes: datetime64[ns](3), float64(5), int64(6), object(10)
memory usage: 237.4+ MB
```

```
In [11]: #Change data type
# Credit card should be integer

df_train.credit_card_number = df_train.credit_card_number.astype('category')
```

```
df_train.is_fraud = df_train.is_fraud.astype('category')
df_train.hour_of_day = df_train.hour_of_day.astype('category')
```

In [12]: df\_train.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1296675 entries, 0 to 1296674
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   transaction_time      1296675 non-null  datetime64[ns]
1   credit_card_number    1296675 non-null  category
2   merchant              1296675 non-null  object
3   category              1296675 non-null  object
4   amount(usd)           1296675 non-null  float64
5   first                 1296675 non-null  object
6   last                  1296675 non-null  object
7   gender                1296675 non-null  object
8   street                1296675 non-null  object
9   city                  1296675 non-null  object
10  state                 1296675 non-null  object
11  zip                   1296675 non-null  int64
12  lat                   1296675 non-null  float64
13  long                  1296675 non-null  float64
14  city_pop              1296675 non-null  int64
15  job                   1296675 non-null  object
16  dob                   1296675 non-null  datetime64[ns]
17  transaction_id        1296675 non-null  object
18  unix_time             1296675 non-null  int64
19  merch_lat             1296675 non-null  float64
20  merch_long            1296675 non-null  float64
21  is_fraud              1296675 non-null  category
22  time                  1296675 non-null  datetime64[ns]
23  hour_of_day           1296675 non-null  category
dtypes: category(3), datetime64[ns](3), float64(5), int64(3), object(10)
memory usage: 212.7+ MB
```

## EDA

In [13]: np.round(df\_train.describe(), 2)

Out[13]:

	amount(usd)	zip	lat	long	city_pop	unix_time	merch_lat	merch_long
<b>count</b>	1296675.00	1296675.00	1296675.00	1296675.00	1296675.00	1.296675e+06	1296675.00	1296675.00
<b>mean</b>	70.35	48800.67	38.54	-90.23	88824.44	1.349244e+09	38.54	-90.23
<b>std</b>	160.32	26893.22	5.08	13.76	301956.36	1.284128e+07	5.11	13.77
<b>min</b>	1.00	1257.00	20.03	-165.67	23.00	1.325376e+09	19.03	-166.67
<b>25%</b>	9.65	26237.00	34.62	-96.80	743.00	1.338751e+09	34.73	-96.90
<b>50%</b>	47.52	48174.00	39.35	-87.48	2456.00	1.349250e+09	39.37	-87.44
<b>75%</b>	83.14	72042.00	41.94	-80.16	20328.00	1.359385e+09	41.96	-80.24
<b>max</b>	28948.90	99783.00	66.69	-67.95	2906700.00	1.371817e+09	67.51	-66.95

In [14]: groups = [pd.Grouper(key="transaction\_time", freq="1W"), "is\_fraud"]  
df\_ = df\_train.groupby(by=groups).agg({"amount(usd)": "mean", "transaction\_id": "count"}).reset\_index()

In [15]: df\_

Out[15]:

	transaction_time	is_fraud	amount(usd)	transaction_id
0	2019-01-06	0	65.783042	9633
1	2019-01-06	1	548.677660	47
2	2019-01-13	0	68.687553	11896
3	2019-01-13	1	542.686240	125
4	2019-01-20	0	67.190627	11777
...	...	...	...	...
149	2020-06-07	1	670.901111	81
150	2020-06-14	0	69.000045	19835
151	2020-06-14	1	571.120225	89
152	2020-06-21	0	69.842152	17778
153	2020-06-21	1	457.665366	164

154 rows × 4 columns

In [16]: `df_train.shape`

Out[16]: (1296675, 24)

```
In [17]: def add_traces(df, x, y, hue, mode, cmap, showlegend=None):
    name_map = {1:"Yes", 0:"No"}
    traces = []
    for flag in df[hue].unique():
        traces.append(
            go.Scatter(
                x = df[df[hue]==flag][x],
                y = df[df[hue]==flag][y],
                mode=mode,
                marker=dict(color=cmap[flag]),
                showlegend=showlegend,
                name=name_map[flag]
            )
        )
    return traces
```

```
In [18]: fig = make_subplots(rows=2, cols=2,
                             specs=[
                                 [{}], [{}],
                                 [{"colspan":2}], None
                             ],
                             subplot_titles=("Amount(usd) over time", "Number of transactions overtime",
                                              "Number of transaction by amount(usd)"))

ntraces = add_traces(df=df_, x='transaction_time', y='amount(usd)', hue='is_fraud', mode='lines',
                     showlegend=True, cmap=['#61E50F', '#D93C1D'])

for trace in ntraces:
    fig.add_trace(
        trace,
        row=1, col=1
    )

ntraces = add_traces(df=df_, x='transaction_time', y='transaction_id', hue='is_fraud', mode='lines',
                     showlegend=False, cmap=['#61E50F', '#D93C1D'])
for trace in ntraces:
    fig.add_trace(
        trace,
        row=1, col=2
    )

ntraces = add_traces(df=df_, x='transaction_id', y='amount(usd)', hue='is_fraud', mode='markers',
                     showlegend=True, cmap=['#61E50F', '#D93C1D'])
```

```


for trace in ntraces:
    fig.add_trace(
        trace,
        row=2,col=1
    )

fig.update_layout(height=780,
                    width=960,
                    legend=dict(title='Is fraud?'),
                    plot_bgcolor='#fafafa',
                    title='Overview'
                    )


fig.show()

```

## Overview



## Number of transaction by amount(usd)



```

In [19]: df_ = df_train.groupby(by=[pd.Grouper(key="transaction_time", freq="1W"),
                                   'is_fraud', 'category']).agg({"amount(usd)": 'mean', "transaction_id": "count"})


fig = px.scatter(df_,
                  x='transaction_time',
                  y='amount(usd)',
                  color='is_fraud',
                  facet_col='category',
                  facet_col_wrap=3,

```

```
        facet_col_spacing=.04,  
        color_discrete_map={0: '#61E50F', 1: '#D93C1D'}  
    )  
  
    fig.update_layout(height=1400,  
                      width=960,  
                      legend=dict(title='Is fraud?'),  
                      plot_bgcolor='#fafafa'  
                      )  
  
    fig.update_yaxes(matches=None)  
    fig.for_each_yaxis(lambda yaxis: yaxis.update(showticklabels=True))  
    fig.for_each_xaxis(lambda xaxis: xaxis.update(showticklabels=True, title=''))  
  
    fig.show();
```








```
In [20]: df_ = df_train.groupby(by=[pd.Grouper(key="transaction_time", freq="1M"),
                                     'is_fraud', 'category']).agg({"amount(usd)": 'sum', "transaction_id": "count"})

fig = px.area(
    df_[df_.is_fraud==1],
    x='transaction_time',
    y='amount(usd)',
    color='category',
    color_discrete_sequence=px.colors.qualitative.Dark24
)

fig.update_layout(height=600,
                  width=960,
                  legend=dict(title='Categories'),
                  plot_bgcolor='#fafafa'
)

fig.show()
```



```
In [21]: # Specified list of 12 merchants with the highest number of transactions.
top12_merchants = df_train.merchant.value_counts()[:12]
```

```
df_ = df_train.groupby(by=[pd.Grouper(key="transaction_time", freq="1W"), 'is_fraud',  
                             'merchant']).agg({"amount(usd)": 'mean', "transaction_id": "count"}).reset_  
  
df_ = df_[df_.merchant.isin(top12_merchants.index)]
```

```
In [22]: fig = px.scatter(df_,  
                        x='transaction_time',  
                        y='amount(usd)',  
                        color='is_fraud',  
                        facet_col='merchant',  
                        facet_col_wrap=3,  
                        facet_col_spacing=.06,  
                        category_orders={'merchant': top12_merchants.index}, # order the subplots  
                        color_discrete_map={1: '#61E50F', 0: '#D93C1D'})  
  
fig.update_layout(height=1200,  
                  width=960,  
                  title='Top 12 merchants with highest number of transactions per week',  
                  legend=dict(title='Is fraud?'),  
                  plot_bgcolor='#fafafa')  
  
fig.update_yaxes(matches=None)  
fig.for_each_yaxis(lambda yaxis: yaxis.update(showticklabels=True))  
fig.for_each_xaxis(lambda xaxis: xaxis.update(showticklabels=True, title=''))  
  
fig.show()
```

Top 12 merchants with highest number of transactions per week



```
In [23]: groups = ['is_fraud', 'job']
df_ = df_train.groupby(by=groups).agg({"amount(USD)": 'mean', "transaction_id": "count"}).fillna(0).reset_index()

# Top 10 jobs had most fraud transactions.
df_ = df_[df_.is_fraud==1].sort_values(by='transaction_id',
                                       ascending=False).drop_duplicates('job', keep='first').iloc[:10]
df_
```

```
Out[23]:
```

	is_fraud	job	amount(USD)	transaction_id
780	1	Materials engineer	561.092097	62
971	1	Trading standards officer	478.137143	56
802	1	Naval architect	653.563962	53
681	1	Exhibition designer	524.067255	51
933	1	Surveyor, land/geomatics	510.914800	50
781	1	Mechanical engineer	531.585714	49
845	1	Prison officer	453.897500	48
875	1	Quantity surveyor	591.754167	48
536	1	Audiological scientist	660.311739	46
604	1	Copywriter, advertising	458.743556	45

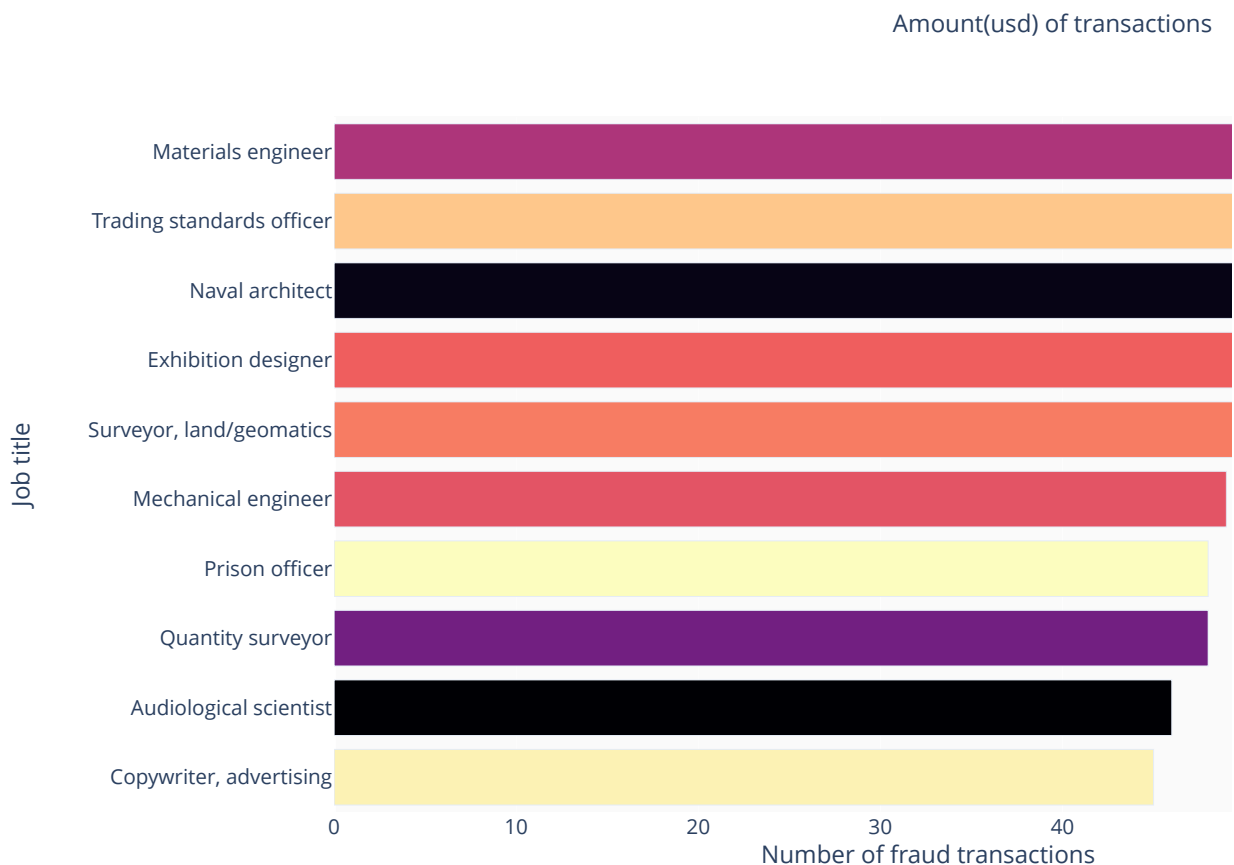
```
In [24]: fig = px.bar(df_,
                    y='job', x='transaction_id',
                    color='amount(USD)',
                    color_continuous_scale=px.colors.sequential.Magma,
                    labels={'job': 'Job title',
                           'transaction_id': 'Number of fraud transactions'},
                    category_orders = {"job": df_.job.values},
                    width=960,
                    height=600)

fig.update_layout(
    title=dict(
        text='Amount(USD) among top 10 jobs with the most fraud transactions'
    ),
    plot_bgcolor='#fafafa'
)

fig.update_coloraxes(
    colorbar=dict(
        title='Amount(USD) of transactions',
        orientation='h',
        x=1
    ),
    reversescale=True
)

fig.show()
```

## Amount(usd) among top 10 jobs with the most fraud transactions



In [25]: *#Which credit card number/ credit card holder has most fraud transaction?*

```
groups = ['credit_card_number']
df_ = df_train.groupby(by=groups).agg({"amount(usd)":'mean', "transaction_id":'count'}).fillna(0).reset_index()
df_.sort_values('transaction_id', ascending=False, inplace=True)
df_ = df_.head(10)
```


In [26]: df\_ = df\_train[df\_train.is\_fraud==1].groupby(by='hour\_of\_day').agg({'transaction\_id':'count'}).reset\_index()

```
fig = px.bar(data_frame=df_,
             x='hour_of_day',
             y='transaction_id',
             labels={'transaction_id':'Number of transaction'})

fig.update_layout(
    title=dict(
        text='Number of FRAUD transactions by hours of day'
    ),
    plot_bgcolor='#fafafa'
)

fig.update_xaxes(type='category')
```

## Number of FRAUD transactions by hours of day



In [27]: `df_train.dtypes`

```
Out[27]: transaction_time    datetime64[ns]
credit_card_number         category
merchant                   object
category                   object
amount(usd)                float64
first                      object
last                      object
gender                     object
street                     object
city                       object
state                      object
zip                        int64
lat                        float64
long                       float64
city_pop                   int64
job                        object
dob                       datetime64[ns]
transaction_id             object
unix_time                  int64
merch_lat                  float64
merch_long                 float64
is_fraud                   category
time                       datetime64[ns]
hour_of_day                category
dtype: object
```

```
In [28]: # Identify non-numeric columns
non_numeric_cols = df_train.select_dtypes(exclude=np.number).columns
```

In [29]: `non_numeric_cols`

```
Out[29]: Index(['transaction_time', 'credit_card_number', 'merchant', 'category',
               'first', 'last', 'gender', 'street', 'city', 'state', 'job', 'dob',
               'transaction_id', 'is_fraud', 'time', 'hour_of_day'],
              dtype='object')
```

```
In [30]: columns_to_drop = ['transaction_time', 'credit_card_number', 'merchant', 'category', 'first', 'last', 'gender', 'street', 'city', 'state', 'job', 'dob', 'transaction_id', 'is_fraud', 'time', 'hour_of_day']
```


```
new_df_train = df_train.drop(columns=columns_to_drop)
```

```
In [31]: new_df_train.corr()
```

```
Out[31]:
```

	amount(usd)	zip	lat	long	city_pop	unix_time	merch_lat	merch_long
amount(usd)	1.000000	0.001843	-0.001926	-0.000187	0.005818	-0.000293	-0.001873	-0.000151
zip	0.001843	1.000000	-0.114290	-0.909732	0.078467	0.000670	-0.113561	-0.908924
lat	-0.001926	-0.114290	1.000000	-0.015533	-0.155730	0.000632	0.993592	-0.015509
long	-0.000187	-0.909732	-0.015533	1.000000	-0.052715	-0.000642	-0.015452	0.999120
city_pop	0.005818	0.078467	-0.155730	-0.052715	1.000000	-0.001714	-0.154781	-0.052687
unix_time	-0.000293	0.000670	0.000632	-0.000642	-0.001714	1.000000	0.000561	-0.000635
merch_lat	-0.001873	-0.113561	0.993592	-0.015452	-0.154781	0.000561	1.000000	-0.015431
merch_long	-0.000151	-0.908924	-0.015509	0.999120	-0.052687	-0.000635	-0.015431	1.000000

```
In [32]: sns.heatmap(new_df_train.corr(), cmap="Greens", annot=True)
plt.show()
```



```
In [33]: features = ['transaction_id', 'hour_of_day', 'category', 'amount(usd)', 'merchant', 'job']

x = df_train[features].set_index("transaction_id")
y = df_train['is_fraud']

print('X shhape:{}\ny shape:{}'.format(x.shape, y.shape))

X shhape:(1296675, 5)
y shape:(1296675,)
```

```
In [34]: from sklearn.preprocessing import OrdinalEncoder

enc = OrdinalEncoder(dtype=np.int64)
enc.fit(x.loc[:, ['category', 'merchant', 'job']])
```



```
x.loc[:, ['category', 'merchant', 'job']] = enc.transform(x[['category', 'merchant', 'job']])
```

```
In [35]: x[['category', 'merchant', 'job' ]]
```

```
Out[35]:
```

	category	merchant	job
transaction_id			
0b242abb623afc578575680df30655b9	8	514	370
1f76529f8574734946361c461b024d99	4	241	428
a1a22d70485983eac12b5b88dad1cf95	0	390	307
6b849c168bdad6f867558c3793159a81	2	360	328
a41d7549acf90789359a9aa5346dcb46	9	297	116
...	...	...	...
440b587732da4dc1a6395aba5fb41669	0	499	215
278000d2e0d2277d1de2f890067dcc0a	1	2	360
483f52fe67fabef353d552c1e662974c	1	599	308
d667cdcbadaaed3da3f4020e83591c83	1	509	485
8f7c8e4ab7f25875d753b422917c98c9	1	370	467

1296675 rows × 3 columns

```
In [36]: from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
```

```
def select_features(x_train, y_train, x_test):
    fs = SelectKBest(score_func=chi2, k='all')
    fs.fit(x_train, y_train)
    x_train_fs = fs.transform(x_train)
    x_test_fs = fs.trannsform(x_test)
    return X_train_fs, X_test_fs, fs
```

```
In [37]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, stratify=y)
print('X_train shape:{}\ny_train shape:{}'.format(x_train.shape,y_train.shape))
print('X_test shape:{}\ny_test shape:{}'.format(y_test.shape,y_test.shape))
```

```
X_train shape:(1037340, 5)
y_train shape:(1037340,)
X_test shape:(259335,)
y_test shape:(259335,)
```


```
In [38]: from sklearn.tree import DecisionTreeClassifier

dcstree = DecisionTreeClassifier(random_state=42)
dcstree.fit(x_train, y_train)

y_pred = dcstree.predict(x_test)
```

```
In [39]: fig = plt.figure(figsize=(8,6))
cfs_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(cfs_matrix, cmap='viridis', annot=True, fmt='d', annot_kws=dict(fontsize=14))
```

```
Out[39]: <AxesSubplot:>
```



With Decision Tree we have F1-Score = 0.69 for label 1

## SMOTE

In [40]: `pip install imblearn`

```
Collecting imblearn
  Downloading imblearn-0.0-py2.py3-none-any.whl (1.9 kB)
Collecting imbalanced-learn
  Downloading imbalanced_learn-0.12.0-py3-none-any.whl (257 kB)
----- 257.7/257.7 kB 659.8 kB/s eta 0:00:00
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\lenovo\anaconda3\lib\site-packages
(from imbalanced-learn->imblearn) (2.2.0)
Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\lenovo\anaconda3\lib\site-packages
(from imbalanced-learn->imblearn) (1.0.2)
Requirement already satisfied: numpy>=1.17.3 in c:\users\lenovo\anaconda3\lib\site-packages (from
imbalanced-learn->imblearn) (1.21.5)
Requirement already satisfied: scipy>=1.5.0 in c:\users\lenovo\anaconda3\lib\site-packages (from i
mbalanced-learn->imblearn) (1.9.1)
Collecting joblib>=1.1.1
  Using cached joblib-1.3.2-py3-none-any.whl (302 kB)
Installing collected packages: joblib, imbalanced-learn, imblearn
  Attempting uninstall: joblib
    Found existing installation: joblib 1.1.0
    Uninstalling joblib-1.1.0:
      Successfully uninstalled joblib-1.1.0
Successfully installed imbalanced-learn-0.12.0 imblearn-0.0 joblib-1.3.2
Note: you may need to restart the kernel to use updated packages.
```

```

WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -rotobuf (c:\users\lenovo\anaconda3\lib\site-packages)

```

```

In [41]: from imblearn.over_sampling import SMOTE
smote = SMOTE(sampling_strategy={1:48050}, random_state=42)

x_train_smote, y_train_smote = smote.fit_resample(x_train.astype('float'), y_train)
print("Before SMOTE:", Counter(y_train))
print("After SMOTE:", Counter(y_train_smote))

```

```

Before SMOTE: Counter({0: 1031335, 1: 6005})
After SMOTE: Counter({0: 1031335, 1: 48050})

```

```

In [42]: class test_model:
    from sklearn.metrics import classification_report
    def __init__(self):
        self.metrics = ['prfs', 'auc', 'acc']

    def fit_predict(model, x_train, x_test, y_train, y_test):
        model = model
        model.fit(x_train, y_train)
        y_pred = model.predict(x_test)
        return y_pred

    def evaluate(y_pred, metrics):
        results = {}
        for metric in metrics:
            if metric == 'prfs':
                prfs = classification_report(y_test, y_pred)
                results['prfs'] = prfs
            elif metric == 'auc':
                auc_score = roc_auc_score(y_test, y_pred)
                results['auc'] = auc_score
            elif metric == 'acc':
                results['acc'] = accuracy_score(y_test, y_pred)
            else:
                print('Not available metric!')
        return results

```

```

In [43]: from sklearn.ensemble import RandomForestClassifier

# Specify your metric here
metrics = ['prfs']
print("=====SMOTE=====")
RDForest_eval = test_model.evaluate(y_pred=test_model.fit_predict(RandomForestClassifier(random_state=42),
                                                                    x_train_smote,
                                                                    x_test,
                                                                    y_train_smote,
                                                                    y_test),
                                   metrics=metrics)

print("Random Forest model evaluate:\n", RDForest_eval['prfs'])

```

```

=====SMOTE=====
Random Forest model evaluate:

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	257834
1	0.79	0.75	0.77	1501
accuracy			1.00	259335
macro avg	0.90	0.88	0.89	259335
weighted avg	1.00	1.00	1.00	259335

With RandomForestClassifier we have better F1-Score = 0.77 for label1

Try tuning some important Hyperparameter for RDF

```
In [44]: # Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 50, stop = 200, num = 4)]
# Number of features to consider at every split
max_features = ['auto', 'sqrt']
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(10, 50, num = 5)]
max_depth.append(None)
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4]

# Create the random grid
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf
              }

print(random_grid)

{'n_estimators': [50, 100, 150, 200], 'max_features': ['auto', 'sqrt'], 'max_depth': [10, 20, 30, 40, 50, None], 'min_samples_split': [2, 5, 10], 'min_samples_leaf': [1, 2, 4]}
```

```
In [45]: from sklearn.metrics import f1_score, make_scorer
f1 = make_scorer(f1_score, greater_is_better=True, pos_label=1)
```

```
In [46]: rf = RandomForestClassifier()
rf_random = RandomizedSearchCV(estimator = rf,
                              param_distributions = random_grid,
                              n_iter = 50,
                              cv = 5,
                              verbose=2,
                              random_state=42
                              )

#Fit and show the best parameters
rf_random.fit(x_train, y_train)
print(rf_random.best_estimator_)
```

[illegible]

[illegible]

```
[CV] END max depth=None, max features=sqrt, min samples leaf=4, min samples split=5, n estimators=
```

[illegible]



25/28

[illegible]

[illegible]

	precision	recall	f1-score	support
0	1.00	1.00	1.00	257834
1	0.87	0.72	0.79	1501
accuracy			1.00	259335
macro avg	0.93	0.86	0.89	259335
weighted avg	1.00	1.00	1.00	259335

After tuning, we have F1-Score = 0.79 for label 1.

```
In [48]: df_test = df_test.rename(columns={"trans_date_trans_time": "transaction_time",
                                         "cc_num": "credit_card_number",
                                         "amt": "amount(USD)",
                                         "trans_num": "transaction_id"})

# Apply function utcfromtimestamp and drop column unix_time
df_test['time'] = df_test['unix_time'].apply(datetime.utcfromtimestamp)

# Add column hour of day
df_test['hour_of_day'] = df_test.time.dt.hour

df_test = df_test[features].set_index("transaction_id")
enc = OrdinalEncoder(dtype=np.int64)
enc.fit(df_test.loc[:, ['category', 'merchant', 'job']])

df_test.loc[:, ['category', 'merchant', 'job']] = enc.transform(df_test[['category', 'merchant', 'job']])
```

```
In [49]: y_pred = rf_random.predict(df_test)
y_proba = rf_random.predict_proba(df_test)[:, 1]
```

```
In [50]: df_test["Fraud_Proba"] = y_proba
df_test["Fraud_Predict"] = y_pred
```

```
In [51]: result = df_test[["Fraud_Proba", "Fraud_Predict"]]
# Store result in a CSV file
result.to_csv(r"./PredictFraud_Result.csv")
```

```
In [52]: result
```

```
Out[52]:
```

	Fraud_Proba	Fraud_Predict
transaction_id		
2da90c7d74bd46a0caf3777415b3ebd3	0.00	0
324cc204407e99f51b0d6ca0055005e7	0.00	0
c81755dbbba9d5c77f094348a7579be	0.00	0
2159175b9efe66dc301f149d3d5abf8c	0.00	0
57ff021bd3f328f8738bb535c302a31b	0.00	0
...	...	...
9b1f753c79894c9f4b71f04581835ada	0.00	0
2090647dac2c89a1d86c514c427f5b91	0.00	0
6c5b7c8add471975aa0fec023b2e8408	0.00	0
14392d723bb7737606b2700ac791b7aa	0.00	0
1765bb45b3aa3224b4cdcb6e7a96cee3	0.01	0

555719 rows × 2 columns

```
In [ ]:
```